

Warm-up Auto-send Plan — Reallocating Human Attention in B2B Outreach

Status: SHIPPED — PR0–PR3 merged 2026-06-05; system armed **Date:** 2026-06-05 **Owner:** Gary Teh (human) + Claude Code (AI assist) **Repos affected:** [TrueSightDAO/go_to_market](#) (implementation), [TrueSightDAO/agentive_ai_context](#) (this plan) **Related docs:** [HIT_LIST_STATE_MACHINE.md](#) (the 13-state pipeline this modifies), [PARTNER_OUTREACH_PROTOCOL.md](#) §9 (the human-in-the-loop convention this graduates from), [EMAIL360.md](#) (the DTC retention-loop sibling — this plan covers the **B2B outbound** loop)

TL;DR (for the Beer Hall)

We studied 60 days of Agroverse retail-partner outreach data and found the human bottleneck is in the wrong place. Gary spends his outreach time clicking “send” on AI-drafted first-touch emails — 98 of which are currently waiting a **median 24 days** for that click — while the rare prospects who actually *reply* wait **29 hours median (worst case 10 days)** for his response. The replies are where partners come from: all 14 current retail partners trace back to a human conversation, not a first-touch email.

The change: the AI now sends the routine, machine-checked first-touch emails itself (drip-fed, capped per day, kill-switched), and Gary’s attention is redirected to the moment a human replies — every reply gets an AI-drafted response within the hour for him to edit and approve. Nothing a prospect writes is ever answered by an unattended machine.

Why DAO members should care: same human hours → more genuine partner conversations per week → faster Agroverse retail expansion → more cacao moving → more trees planted. And the playbook is reusable for any DAO outreach surface.

1. Assessment — what the data says

1.1 Why now

Three artifacts surfaced on 2026-06-05 that prompted the audit: two AI-automated cold-pitch email sequences received at DAO inboxes (burner domains, rotating fake personas, a reply-handling bot that answered a buying signal with the single word “No”), and a Stanford GSB clip of Ken Griffin arguing AI has erased the outbound-scale moat — his example being a 25-year-old who used signal-triggered, AI-personalized outreach to build a pet-insurance business that sold for \$1B. The spam specimens demonstrate the failure mode: **they automated the cheap part (sending) and abandoned the prospect at the expensive part (the reply)**. The Griffin clip demonstrates the prize: the differentiator is signal quality and reply handling, not send volume. Our pipeline had the inverse problem — excellent reply handling (when it happens) but a human gate throttling the cheap part.

1.2 Pipeline state (Hit List workbook, 2026-06-05)

Metric	Value
Hit List rows total	669 across 14 states
Rows in AI: Warm up prospect (ready for first touch)	79
Rows in Manager Follow-up	34

Metric	Value
Partnered	14
Drafts generated (Email Agent Drafts, all time)	470 — 183 sent / 189 discarded / 98 pending_review
Pending-review backlog age	median 24 days (74 warm-ups, 24 follow-ups)
Send pattern	bursty — 93 sends on Sun 2026-05-03, 74 on Fri 2026-04-24 (weekend batch sessions)
Draft→send latency (sent cohort)	median 1.1 d, p75 = 12 d , max 30 d

1.3 Engagement ground truth (Gmail thread audit, 232 sent threads)

Metric	Value
Genuine prospect replies (tracking-tool noise removed)	14 ≈ 6% — healthy for cold B2B
Replies that were real opportunities	3 — one shop asked for samples (shipped), one live wholesale conversation, one became a Partnered store
Gary's response latency to genuine replies	median 28.8 h, max 239 h
Response latency on the thread that converted to Partnered	0.1–1 h
Copy quality complaints in any reply	0 (worst feedback: a constructive “show product photos on your site”)

The pixel-based “97% open rate” is discounted — Apple Mail / Gmail proxy prefetch fires tracking pixels without a human open. Replies are the only ground truth.

1.4 Verdict

1. **The send-click is the pipeline bottleneck.** Cron jobs qualify, enrich, and draft 24/7; everything then queues on a weekend batch session.
2. **Human attention is mis-allocated.** ~94% of sends never get a reply — yet they consume the operator's review time, while the 6% who raise their hand wait days. The fastest-handled reply became a partner; the slowest soft-no sat 10 days.
3. **The graduation criterion was already met.** `HIT_LIST_STATE_MACHINE.md` (operator review loop) defined the test: *“If reply rate stays healthy AND no pattern of bad copy in the lint-reviewed sample, that's the signal to drop the unflagged tier to fully auto-send.”* Reply rate ~6%, bad-copy complaints zero. The evidence gate is cleared.
4. **Burst sending is itself a risk.** 93 sends in one day from one Gmail account is a deliverability gamble; a daily drip is safer and produces a steady, manageable reply flow instead of reply pile-ups.

Principle adopted: allocate human attention by demonstrated prospect intent, not by pipeline stage.

2. Design principles and guardrails

1. **Auto-send only the machine-checked tier.** A draft auto-sends only if the 12-rule linter (`preview_warmup_drafts.py`) raises **zero red and zero yellow flags**. Flagged drafts stay in the human review queue exactly as today.

2. **Replies are always human-approved.** Automation may *draft* a response and *notify*, never send one. (Lesson from the spam specimens: the moment of engagement is precisely where unattended automation destroys deals.)
 3. **Drip, don't burst.** Default cap 12 sends/day, weekdays, business hours (Pacific). Protects sender reputation and smooths reply load.
 4. **Kill switch + dry-run default.** Repo variable `WARMUP_AUTOSEND_ENABLED` gates the workflow; the script defaults to `--dry-run` and only sends with explicit `--execute`.
 5. **Full audit trail.** Every auto-send flows through the existing `sync_email_agent_followup.py` reconcile (status flip, label swap `AI/Warm-up` → `AI/Sent Warm-up`, `Email Agent Follow Up` log row), so the sent history is indistinguishable in structure from human sends.
 6. **Anti-patterns we will never adopt** (from the specimens): rotating fake personas, lookalike burner domains, contradictory canned answers, abandoning a thread after a warm reply, volume masquerading as qualification.
 7. **First touches and cadenced no-reply follow-ups only.** Anything after a prospect speaks is human territory.
-

3. Implementation plan — sequenced PRs on `go_to_market`

PR1 — Auto-send the linter-clean warm-up tier

New `scripts/send_clean_warmup_drafts.py`:

- Reads `Email Agent Drafts` rows `status='pending_review'`, `gmail_label='AI/Warm-up'` (same selection as `preview_warmup_drafts.py`).
- Reuses the existing `_lint()` rules and adds one new **red** rule: `email_domain_mismatch` — recipient is on a custom (non-freemail) domain that differs from the Hit List `Website` domain. (Catches the observed misfire where a warm-up went to a shop's *marketing agency* inbox.)
- Clean tier = zero red **and** zero yellow flags. Everything else is skipped and left for the human review loop.
- Sends via Gmail API `drafts().send()` (the draft already exists in the operator's mailbox — sending preserves thread, attachments, and labels).
- Caps sends per run (`--max-sends`, default 12), oldest first; optionally `--require-hosts-circles no` to exclude high-leverage `AW=Yes` prospects from auto-send (default: `AW=Yes` drafts are **excluded** — they get human eyes, consistent with the linter's blue flag).
- After sending, invokes the `sync_email_agent_followup.py` reconcile path so status/labels/log update in the same run.
- `--dry-run` default prints the would-send list with lint results.

New workflow `.github/workflows/warmup_autosend.yml`:

- `schedule`: weekdays 17:05 UTC (≈ 9 am PT) + `workflow_dispatch`.
- Gated on repo variable `WARMUP_AUTOSEND_ENABLED == 'true'`.
- Secrets: `GOOGLE_CREDENTIALS_JSON`, `GMAIL_TOKEN_JSON` (both already exist on the repo for sibling workflows).
- `concurrency` group to prevent overlapping runs.

PR2 — Reply acceleration: auto-draft + notify

New `scripts/draft_prospect_reply_responses.py`:

- Detects genuine inbound replies on warm-up / follow-up threads (reuses the detection logic in `backfill_warmup_reply_remarks.py`; filters tracking-tool noise such as Mailsuite reminder messages).
- For each new reply with no existing response draft: builds full context (thread history, Hit List row, DApp Remarks) and generates a Grok-drafted reply **in-thread** as a Gmail draft, labeled `AI/Reply-Draft`.
- Sends a notification email to the operator (subject: `[Reply] <shop> – <one-line gist>`; body: the prospect's message, the draft text, a Gmail deep-link) so the median human response drops from ~29 h to same-session.
- Idempotent on Gmail `message_id` (a processed-marker via the existing DApp Remarks audit pattern).

Workflow `reply_draft_responses.yml`: hourly, same secrets + `GROK_API_KEY`.

PR3 — Reply classification + soft-no auto-parking

New `scripts/classify_warmup_replies.py`:

- LLM-classifies each genuine inbound reply into:
`interested` | `question` | `soft_no_with_date` | `soft_no` | `hard_no` | `wrong_contact` | `referral`.
- Automatic Hit List updates for the unambiguous negatives:
- `soft_no_with_date` (“reach out in September”) → Status `Deferred` / `Revisit later` + `Follow Up Date` set from the parsed date.
- `hard_no` (“we are not interested”) → Status `Rejected`.
- `referral` (named alternative shops) → referred names appended to `Sales Process Notes` for discovery triage.
- `interested` / `question` / ambiguous → untouched (they ride PR2's notification path).
- Every classification writes a `DApp Remarks` audit row with the model rationale, using the shared `hit_list_dapp_remarks_sheet.py` apply-semantics.

Runs inside PR2's hourly workflow (classification informs the notification subject line).

4. Pre-flight checklist

- [x] Gmail user OAuth token valid with `gmail.modify` (verified 2026-06-05)
 - [x] Linter + clean-tier definition exists (`preview_warmup_drafts.py`)
 - [x] `reconcile_drafts_status_to_sent()` exists (`sync_email_agent_followup.py`)
 - [x] Evidence gate met (reply rate ~6%, zero copy complaints — §1.4)
 - [x] Repo secrets present on `go_to_market`: `G00GLE_CREDENTIALS_JSON`, `GMAIL_TOKEN_JSON`, `GROK_API_KEY` (used by sibling workflows)
 - [x] Repo variable `WARMUP_AUTOSEND_ENABLED=true` created 2026-06-05
 - [x] Daily cap **12** + `AW=Yes` exclusion confirmed (operator green-lit full execution)
 - [x] First armed run dispatched 2026-06-05 (`workflow_dispatch`, `max_sends=12`)
-

5. Execution roadmap — resume tracker

RESUME HERE → all build units shipped 2026-06-05. Remaining: observe the backlog drain over the coming week and read out the §6 metrics at +30 days.

Unit	Scope	Repo	Merged	Contribution reported
PR0	This plan (md + pdf)	agentic_ai_context	✅ #292	✅ consolidated
PR1	send_clean_warmup_drafts.py + warmup_autosend.yml + email_domain_mismatch lint rule	go_to_market	✅ #156	✅ consolidated
PR2	draft_prospect_reply_responses.py + notification + hourly workflow	go_to_market	✅ #157	✅ consolidated
PR3	classify_warmup_replies.py + soft-no auto-parking	go_to_market	✅ #158	✅ consolidated
Post	Arm WARMUP_AUTOSEND_ENABLED, observe first run, drain 98-draft backlog	operations	✅ armed 2026-06-05, first run dispatched	—

6. Success metrics and rollback

Metric	Baseline (2026-06-05)	Target (30 days)
Pending-review backlog age (median)	24 days	< 2 days
Draft→send latency (p75)	12 days	< 3 days
Operator response latency to genuine replies (median)	28.8 h	< 4 h
Genuine reply rate	~6%	≥ 5% (must not degrade)
Bounce / spam-complaint rate	~0	0 sustained
Genuine partner conversations per week	~1	3+

Rollback: set `WARMUP_AUTOSEND_ENABLED=false` (instant, no deploy); drafts simply accumulate in the human review queue as they do today. PR2/PR3 are read-and-draft only — disabling their workflow restores the status quo.

Watch item: if the auto-sent cohort's reply rate falls materially below the human-reviewed cohort's, re-tighten the clean-tier definition (e.g. require `Hosts Circles` crawl freshness, or re-include yellow flags as blockers) before resuming.

Changelog

- **2026-06-05** — Plan created from the outreach-data audit (470 drafts, 232 thread Gmail audit). Approved by Gary same day; execution started.

- **2026-06-05 (later)** — PR0–PR3 all merged same day. Live dry-runs during build: linter caught `compliance@4life.com` + `filler@godaddy.com` before they could auto-send; classifier matched the manual reading on all 7 real reply threads and refused to clobber 3 operator-triaged rows. System armed (`WARMUP_AUTOSEND_ENABLED=true`), first run dispatched.